

A Survey of Web Testing Frameworks and a Workflow for Testing Javascript Applications

By Taggart Ashby

Computer Science Department, California Polytechnic State University, San Luis Obispo

May 15, 2014

Abstract

Web applications are quickly replacing standalone applications. As with any type of applicaiton, these web applications need to be tested to ensure proper functionality and reliability. There have been substantial efforts to create tools to help with the testing of web applicaitons but there is, as of yet, no standard or recommended workflow to insure speed of development and strength of application.

We have brought together a number of testing tools from the web and created what we feel is a fully-featured, easy to use, testing framework and workflow for javascript web applications.

Make this
better

1 Introduction

Since the advent of the internet, the web has undergone an impressive evolution from plain-text webpages to the highly stylized and functional pages that we see today. In the last four to six years there has been an explosion of new web technologies that make applications more functional: HTML5, CSS3, WebGL, Touch APIs, Geo-location, and a host of others. [18] In that same time period, the number of global internet users has grown to somewhere around 2.5 billion people. [18] The combination of these new technologies and the ever-increasing usage of the internet has led to a significant growth in the number and complexity of web applications.

Like any piece of software, these web applications require testing in order to determine their correctness and give their users the best experience possible. Unfortunately, web applications are new and there are no standard practices regarding how to test them. Their dynamic nature and network communication require a different testing toolkit than their predecessors. In addition, web applications are being produced at an incredible rate and so there's a need for rapid testing.

We can see that this is not a solved problem by looking at some of the biggest web application around. A 2011 study showed that YouTube has at least eight errors, Apple has at least seven, Microsoft at least four, and the list goes on. [13] That may seem like a incredibly small amount for such monoliths, and, honestly, it is fairly small, but those are still errors in the software that ought to be fixed.

What this all boils down to is that there needs to be a framework in place for testing these multi-layer, multi-faceted web applications that is easy to use, quick to develop on, and thorough in its coverage.

1.1 What Has Been Done Already?

Web developers are not clueless, they understand that there is a need for testing of web applications. Unfortunately, web development is still a relatively new domain and developers have not come together to work on the problem of testing but rather have come up with their own proprietary solutions in isolation or merely on a per-project basis. Some of these tools have gained traction and become larger, more general, tools.

1.1.1 Why It Is Insufficient

The problem that followed was that many of the tools attempted to remedy the same problems and so a developer would have to decide between a number of similar tools. Having to wade through a sea of similar tools is time-consuming and that can be the death of a cutting-edge application. The more time an application is off the market, the more time competitors have to develop their own versions. Yet another problem is that deciding

What goes here?

Info info info

between similar tools can be discouraging to the developer that just wants to write their application and test it quickly.

Info info
info

1.2 Problems Still Faced/Why It's So Hard

Testing web applications is difficult. The web offers an impressive set of features, but with those features come some draw-backs. Two of the most frustrating drawbacks for developers are the number of platforms and the web's asynchronous nature.

In the very beginning, web applications were easy to write and test because the only code was straight HTML with a common set of HTML tags. It didn't matter what browser was used to view the content because the viewer was always a home computer rendering plain HTML. As time progressed individual browsers began offering additional features and web applications became more rich but harder to test because the developer would need to run their code on a couple different browsers to make sure they all behaved similarly. In today's technological world, it is nearly impossible to run one's web application on all of the possible browser versions, configurations, and platforms that are available. Web applications are on home computers, phones, tablets, MP3 players, refrigerators [42], and a variety of other devices of varying sizes and capabilities. As a developer, one wants to target as many of those platforms as possible with a single code-base to save oneself time and effort. A naive developer might just release an application on all platforms and hope for the best, but releasing on all platforms without testing will likely drive consumers away if the application fails on their given platform.

The addition of Javascript to web applications led to a golden age of feature-rich applications with exciting UI that had not before been seen. This set of features and UI improvements came with an, albeit necessary, complication of asynchronous execution. Unlike traditional multi-threaded software, however, the developer is not given access to individual threads but is stuck with just the main thread. This sort of implementation is essential for dealing with asynchronous events like button clicks, data loading that doesn't hold up the web page, and a variety of other niceties, but is the cause of major headaches in testing. Without access to individual threads, the developer has to either write synchronous code that will be slow and not user-friendly, or rely on other mechanisms that allow code to be sequenced. Often it is hard to get code properly sequenced and it requires changing straightforward functions to a hideous amalgamation of callbacks.

There is no overarching solution to these problems, but in order to understand the motivations behind this thesis, one must understand the problems facing web application developers.

1.3 Contributions of This Thesis

The goal of this thesis is to bring together tools that are already available for testing and combine them into a single framework that offers web application developers a toolkit for their testing needs. There are plenty of strong tools available and we hope to gather a subset of those tools to cover as much of the testing spectrum as possible.

In addition to creating a single framework, we've set out to weigh the options between the most popular frameworks so that a developer need not just blindly trust our recommendations. Since every project is unique, a single set of tools cannot possibly cover the gamut. In addition to the single set of recommendations, the most popular tools are discussed so that the developer is free to pick and choose as they see fit.

2 Background

In an effort to make the rest of the paper more understandable and to define the terms used in the remainder of the paper, we will outline the various concepts and technologies surrounding the web. If one is familiar with HTML5, JavaScript, and Testing concepts and statistics, it will likely be safe to skip this background section.

2.1 Web Application Strengths

Web applications are great for both developers and consumers. They have a number of characteristics that make them more desirable than shrink-wrap software. These attributes are versioning, availability, platform independence, and the ability to rapidly prototype.

2.1.1 Versioning

Web applications can be updated at any time with instant roll-out and feedback. It requires absolutely no user interaction because the web page always serves the latest resources. In comparison to something like an operating system that requires users to manually update, this feature leads to applications that can address problems quickly, roll-out security updates instantly, and incorporate consumer feedback much more quickly than any other type of software. Another perk is that users are not given a choice. This can lead to occasional customer outrage, but more often than not it's a powerful gain for both feature improvement and application security.

2.1.2 Availability

Web application are available anywhere there's internet. No need to install anything on a given device most of the time. Due to the ever-increasing popularity of mobile devices, most applications have both mobile and full versions of their software and so the application is "with" the user all the time.

2.1.3 Multi-platform

Very similar to the availability strength, web applications are platform independent. Whether the user is on their phone, tablet, laptop, or PC, the application only requires a web browser with certain functionality and in this day and age all devices come equipped with browsers that should have no problems.

2.1.4 Rapid Prototyping

Web applications are fantastic for rapid prototyping because of all of the strengths discussed above. An idea can be prototyped and posted on a website quickly and rapid iteration can occur without the user even necessarily noticing. There are a number of web tools that will outline where users are clicking, what features they're using, how long they're spending on a given page or step of a process, and a number of other metrics that can make iterations more directed and effective.

2.2 Web Features

Here we'll discuss some of the prevalent technologies that recently appeared in web applications and development. By no means is this an exhaustive list, just a sampling of the more prominent ones.

2.2.1 HTML5

The latest revision of HTML is HTML5 which first debuted on Firefox back in 2009. [18] HTML5 has not been officially recommended by the W3C (World Wide Web Consortium) who is the official keeper and recommender of web standards, however, there is a plan for final recommendation this year, 2014, called Plan 2014. [11] What this means is that in 2014 the W3C will release a stable HTML5 Recommendation which will make HTML5 an official standard and help universalize all the separate implementations.

HTML5 brings a number of new features including: audio and video tags, the canvas element, drag-and-drop functionality, web storage, offline page support, and a number of other modern web features. A lot of HTML5's features revolve around the inclusion of multimedia in web pages and stronger graphical processing and programming.

As the web is being increasingly used on mobile devices, HTML5 also introduces geolocation functionality. This functionality requires user approval, but once that approval is gained, the web application has access to a rough idea of where the user is located. When used on devices that have GPS, such as phones and tablets, this area is more defined and a stronger lock on the user's position can be obtained. This geolocation functionality has opened the door for brand new types of web applications that can much more closely tie in with what the user is seeing or the area that they are in. For example, ads can be more targeted or an application can display distances to nearby points of interest.

JavaScript, though not the only scripting language with web use, is the primary language used to compliment HTML5 when making web applications. It is an interpreted language that is most often associated with client-side functionality, but in recent times has extended to the server-side. JavaScript is a prototype based object oriented scripting language with dynamic typing and first-class functions. Much of JavaScript's popularity stems from this dynamic nature that allows for quick, and often dirty, implementation of features in an iterative environment.

2.2.1.1 Node.JS As a quick aside, it will be useful to have information about Node.JS as it is used in the sample application and the evaluation applicaiton. Node.JS, often called simply Node, is a JavaScript API for creating and running servers. "Node's goal is to provide an easy way to build scalable network programs." [21] Node was created in 2009 and is sponsored by Joyent. Node does not use thread-based networking, instead opting for a single-threaded event loop and non-blocking I/O. Node allows developers to create and control web servers without the need for external software, such as Apache, commonly found in other sites.

Node is found in a number of popular sites including: PayPal, The New York Times, Yahoo!, Uber, LinkedIn, Microsoft Azure, and a number of others. [21]

2.3 Client/Server Architecture

Web applications have a user-facing client-side and a "rear-facing" server-side.

Here we'll discuss what the client-side is and does, what the server-side is and does, and how they communicate.

2.3.1 What is the Client-side Portion?

First, we'll address what a client is. The simple answer is: The User. The User's phone, computer, tablet, or other device communicating with a given website or application is the client. More specifically, the browser or service being using to communicate with the web is the client.

So now that the client has been clarified, what is the client-side portion of an application? The client-side portion interacts with the browser to request assets from the server and then displaying them in a human-readable way. The client-side also deals with the user interface and interactivity of a web page, e.g., the client-side displays a web form, the client fills in the information and clicks "submit" and the client-side packages up the data and sends it to the server-side.

Client-side programming is done in a language that the browser understands, most frequently Javascript. Though not technically programming languages, client-side programming is also done through HTML and CSS which control the look and feel of web

TIE THIS
CLIEN-
T/SERVER
STUFF
into Web
Applica-
tions, be
specific,
not
generic

IS THIS
THE
RIGHT
TERM...

pages.

2.3.2 What is the Server-side Portion?

Once again, we'll outline a server and then delve into server-side programming. A server is the entity, be it machine, person, or cloud service, that stores the information about a given web page or application. The server stores all assets, HTML pages, images, videos, databases, etc, and provides them upon request to the client. The server also handles all of the processing for interactivity, accounts, and a general number-crunching.

The server-side portion of an application, then, stores all of the assets, accounts, and structure of the entire application. The server-side authenticates users, retrieves data from databases, processes large data, and send information, be it assets or results, back to the client to display to the user.

2.3.3 Client and Server Interactions

The client has limited access to resources on any given web application. This access is restricted by the server. Without such restrictions in place, someone could access user databases, credit card information, order reports, etc. In addition, there would be an overwhelming amount of content for a client to parse through, even if they have no malicious intent.

Clients communicate with the Server through a request-response architecture. Clients send a request to the server, sometimes with information attached, such as a form submittal or login, and the server returns a response with either a success/failure message, or additional data.

2.4 What is testing?

IEEE defines software testing as “the dynamic verification of the behavior of a program on a finite set of test cases, suitably selected from the usually infinite executions domain, against the expected behavior.” [2] Put a bit more simply, software testing is the process of making sure the program does what it is supposed to.

There are a number of different types of tests and testing concept, we will just focus on Unit Testing, Continuous Integration, and Acceptance Testing.

2.4.1 Unit Testing

Unit testing is the testing of specific functions and functionality, or “test units.” Test units are “A set of one or more computer program modules together with associated control data, [...] usage procedures, and operating procedures taht satisfy the following conditions: 1) All modules are from a single computer program. 2) At least one of the new or changed modules in the set has not completed the unit test. 3) The set of modules together with its

associated data and procedures are the sole object of a testing process.” [1] Generally this means the testing of individual functions but can extend to a set of functions that produce a single module of functionality. What sets Unit Testing apart from other types of testing is that tests focus on one portion of the code or system rather than interactions between modules or the entire system.

Here’s a brief example:

```
// Terribly useless code whose purpose is to eventually calculate a magic number
function addTwo (numToAddTo) {
    return numToAddTo + 2;
}

function twoTimes (numToMultiply) {
    var sum = 0;

    for (var i = 0; i < numToMultiply; i++) {
        sum = addTwo(sum);
    }

    return sum;
}

.... // More functions

function calculateMagicNumberPartOne (startingNum) {
    var foo = 3 + addTwo(startingNum);

    return twoTimes(foo);
}

... // More functions

function calculateMagicNumber (startingNum) {
    var numSoFar = calculateMagicNumberPartOne(startingNum);

    numSoFar = calculateMagicNumberPartTwo(numSoFar);
    numSoFar = calculateMagicNumberPartThree(numSoFar);
    ... // More calculateMagicNumberPart(s)

    return numSoFar;
}

function testAddTwo () ... // Unit Test
function testTwoTimes () ... // Unit Test
function testCalculateMagicNumberPartOne () ... // Still a Unit Test, just a
    whole module of functionality
```



```
function testCalculateMagicNumber () ... // NOT a Unit Test because it's testing
    the whole system
```

Unit tests are important because they are done in isolation from one another. The purpose of a unit test is to ensure that once individual modules are put together, the developer need only focus on the interaction between them instead of their individual functionality. Without the assurances of prior unit tests, integration would be a more stressful and mysterious proposition.

2.4.2 Continuous Integration

The original definition of Continuous Integration involves the idea of developers committing from their branches to a master branch daily so that changes are continuously integrated in an effort to make sure all branches can, in fact, be integrated. It has since evolved into a more general idea of a tool that runs commands after each commit to a repository. This commands generally encompass running unit tests, integration tests, user interface tests, and also often publishing and disseminating notifications to interested parties. This is invaluable feedback because the developer can see if they're making progress in the right direction as they write code, as opposed to a situation in which the developer finishes coding and finds out they've made a bad assumption or decision between this commit and the last.

One of the most important aspects of a good piece of CI software is that it is automatic. There should be no need for a developer to do anything besides, perhaps, “turn it on” and start coding. As soon as a tool, CI or otherwise, requires more than one or two interactions to run, it becomes more of a nuisance and less of an asset.

2.4.3 Acceptance Testing

Acceptance testing is testing to confirm that the product meets a certain specification. Often this is making sure that the software meets an SRS (Software Requirements Specification) or a customer's outlined expectations. Often this testing is performed by the user of the software or customer that ordered the software, rather than the developers.

In an effort to automate this, some developers create “User Stories” which are automated tests that run a series of functions described by customers or an SRS. These User Stories are often written in a separate language that is more like english than code. The traditional User Story template is: “As a *role*, I want *goal/desire* so that *benefit*.” For example: “As a supervisor, I want to find my employees by their employee ID.” These user stories are then broken down into more specific steps that can be automated.

2.4.4 Testing Statistics

It is widely accepted that the earlier a developer detects a defect, the easier it is to fix that defect. A famous book, “Code Complete” by Steve McConnell [27], includes the following figure, showing this statistic on a rough graph.

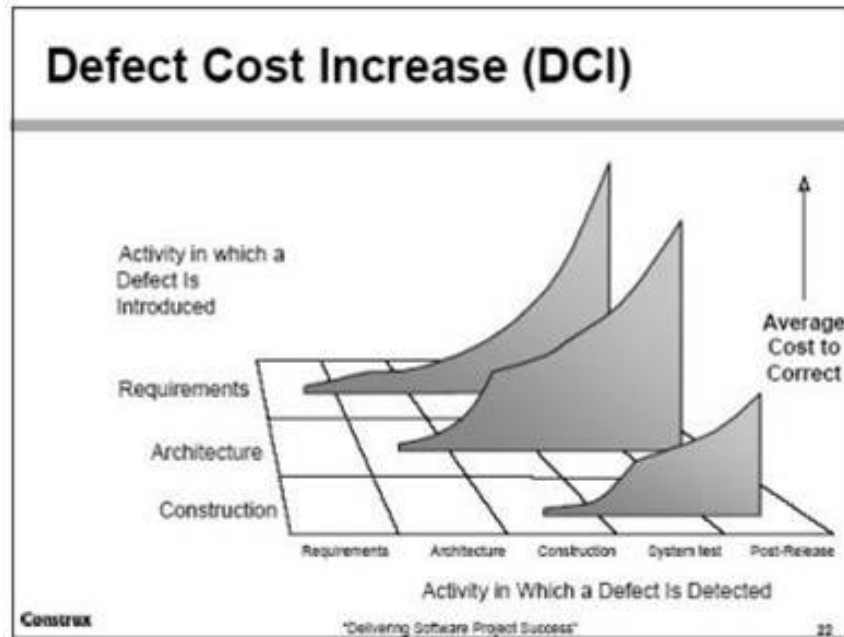


Figure 1: [27]

What we see here is that the earlier a defect is introduced the more it costs to fix it later. Of particular note is that the costs are roughly exponential, not linear. The longer a bug lies undetected, the increasingly harder it becomes to fix that bug. On the flip side, it's relatively easy to fix a bug that was introduced in the same stage that it is found, e.g. introducing a coding bug during “Construction” and fixing it while still writing code, prior to the “System Test” phase.

3 Sample App and Testing Frameworks

3.1 Sample ToDo App

In order to test out the various frameworks that we found online, we needed a small application with some functionality, but without the inherent complexity that was going to come with our evaluation application. After some searching, We decided upon a Todo application made with Express.js, Node.js, and MongoDB. [26] Here is a brief preview of the app:

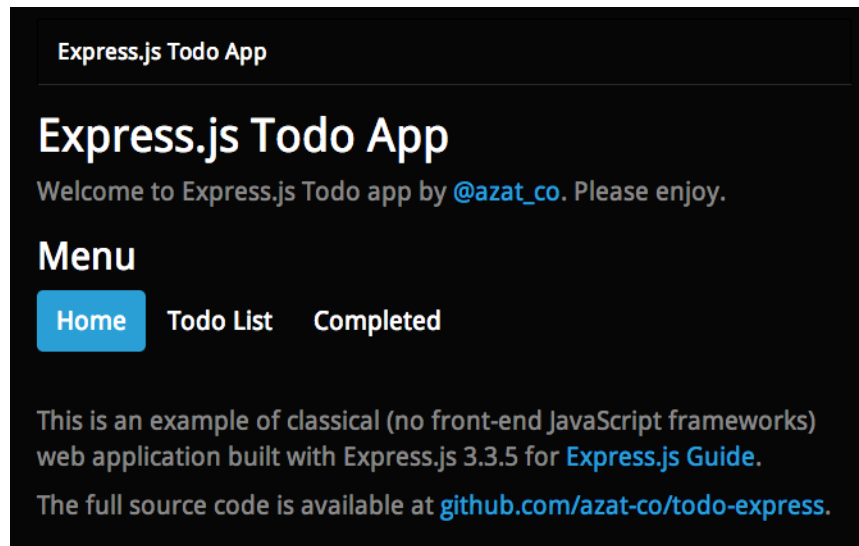


Figure 2: ToDo app home page, simple three-button navigation

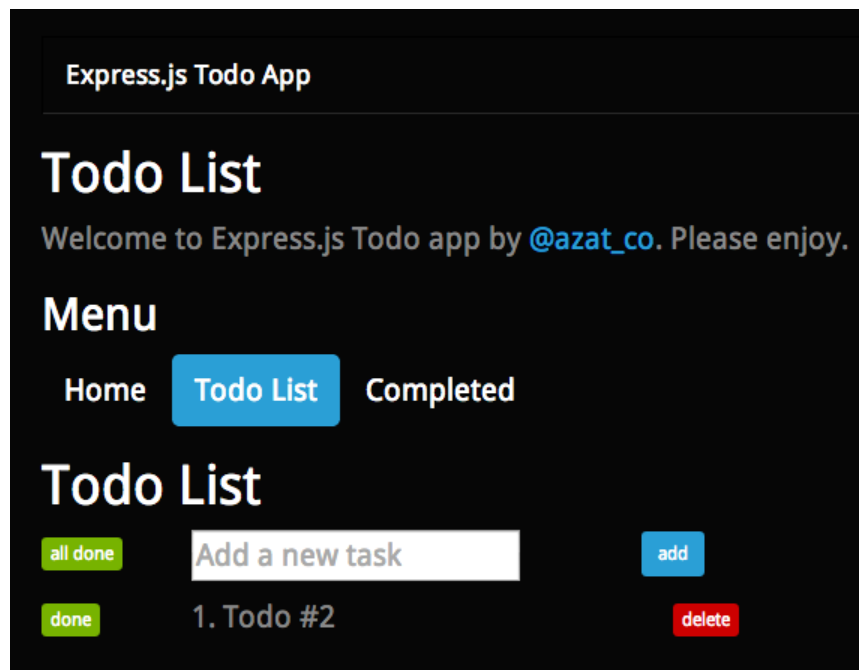


Figure 3: ToDo app task page

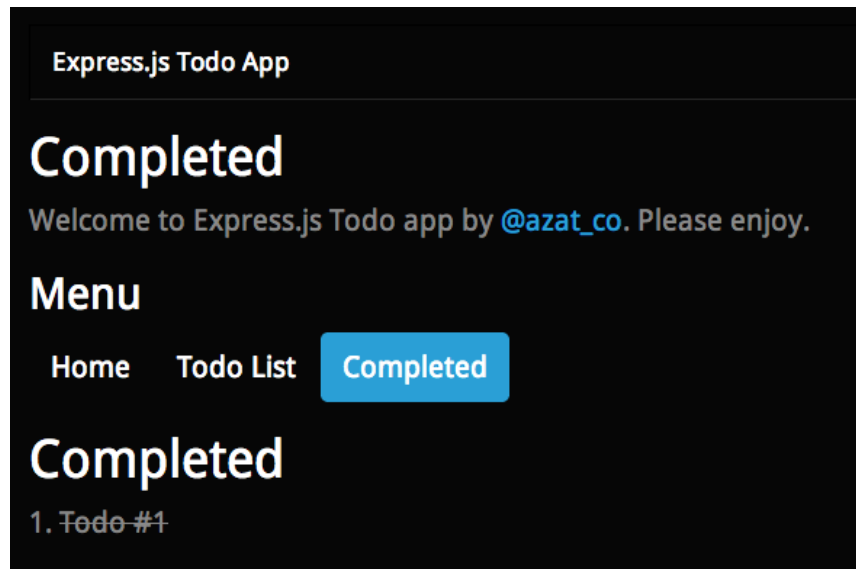


Figure 4: ToDo app completed tasks page

This application was selected because it included Node.js and had enough functionality that tests wouldn't be completely trivial while also not being overwhelming. One of the challenges in testing that we quickly came across was the asynchronous nature of the MongoDB calls that were made in the sample app. Because of this headache, we will be discussing how the following frameworks deal with asynchronous test support as JavaScript is inherently asynchronous.

4 Testing Frameworks

4.0.1 The Average Web Application

In an effort to create a good workflow and collection of tools, we have chosen to focus on web applications that uses Node.js, Javascript, Cloud Services, and a backing Database when searching for our tools. This choice was made because of the popularity of that kind of set up, as well as the fact that our evaluation application was set up in that way.

Although there are multiple languages and tools available for creating a backing server, we chose Node.js because of it's wide use and because it is what we're most familiar with.

Javascript was chosen because it used in over 87% of all websites. [53]

Many web applications use Cloud Services for a number of things, most often they're used for storage of the web content that's being hosted.

Web applications tend to keep track of users which most often constitutes the use of a database behind the scenes to store users and update them as necessary.

THIS IS
LAME

4.0.2 Tool Requirements

Our requirements for the selection of testing tools was relatively simple. In addition to meeting the requirements of our “average web application”, the only other requirement was that the tool was free. What “free” means is that the developer would have access to at least the minimum functionality needed to use the tool on *one* application that is *open-source*. What we quickly found is that web testing tools love open-source projects but are far less willing to provide free functionality to private or closed-source projects.

Occasionally the cost of tools will be discussed but in general we did not review tools that only offered paid options.

Move this
section
elsewhere

4.1 Unit Testing Frameworks

The following frameworks are reviewed in chronological order of our evaluation. As such, comparisons will generally only be drawn between a framework and its predecessors in paper order. After all of the frameworks are outlined and evaluated, a longer section regarding the final decision will be made where comparisons will be drawn between all of them.

An important distinction that needs to be made is that in web applications there is functionality behind the scenes and there is the UI element of the web page. Unit testing frameworks

4.1.1 Jasmine [33]

The first of the unit testing frameworks that we looked at was Jasmine, a Behavior Driven Development (BDD) testing framework. BDD is based on Test Driven Development and adds some simplifications and patterns that attempt to bring both developers and businessmen into the software testing process. The idea behind Behavior Driven Development is that input from non-technical stakeholders as well as technical stakeholders can come together so that everyone understands what the project should do.

This relates to Jasmine is the way that the developer writes Jasmine tests. The idea behind jasmine tests is that they read, as much as possible, like english sentences. As an example, the following is a valid Jasmine test:

```
describe("Arithmetic Test", function () {  
  it("should be able to compute the addition of two numbers", function () {  
    expect(1 + 2).toEqual(3);  
  });  
  
  it("should be able to subtract numbers", function () {  
    var theAnswer = 10;  
  
    expect(12 - 2).toEqual(theAnswer);  
  });  
});
```

```
});  
});
```

The `describe` statement creates a test suite to collect related tests under. Each `it` statement is generally one test with some number of `expect` statements. It may not be particularly pretty, but the above code is quite readable even to someone with no coding experience.

The default Jasmine package had some ineptitudes in dealing with Node.js but there was a package available through npm called `jasmine-node` [15] that gave jasmine access to the node package and functions.

Asynchronous testing using Jasmine was a bit of a chore, but there is support built in. This support comes in the form of a three-set function progression. The first function, `runs(function () { ... })`, runs the asynchronous code. The second function, `waitsFor(function () { ... })`, polls a flag that is set by the first function when it completes. Finally, the third function, another `runs(...)`, houses the assertions to determine if the code ran properly or not. This method requires a fairly substantial amount of typing as well as complicated function chains when more than one asynchronous call is required.

4.1.2 Mocha [52]

Mocha is another BDD testing framework that is almost identical to Jasmine. Unlike Jasmine, however, Mocha is not completely self-contained and requires an additional assertion library of one's choosing. We will outline our choice of Chai.js below.

Mocha, by default, also uses the `describe -> it(\"should...\")` pattern to create tests. Mocha offers a variety of other test interfaces for developer's to write their test code the way they want to in addition to a custom, nearly unstructured variety. A few of the testing interfaces are outlined here:

```
// BDD (default)  
describe('BDD Suite', function () {  
  it('should be test 1', function () {  
    ...  
  });  
});  
  
// TDD  
suite('TDD Suite', function () {  
  test('test 1', function () {  
    ...  
  });  
});  
  
// Exports  
module.exports = {  
  'Exports Suite': {
```

```

    'test 1': function () {
        ...
    }

}

};

// QUnit
suite('QUnit Suite');

test('test 1', function () {
    ...
});

```

The developer doesn't gain anything by using one interface over another except the ability to write tests in a way that makes sense to them. There is no standard for writing assertions, as each assertion library is likely to be slightly different.

Mocha's approach to asynchronous testing is quite a bit more straightforward than Jasmine's. In the `it(...)` anonymous function, there is passed an additional parameter called `done` which is then called when the asynchronous call has finished. This required us to edit the code we were testing to make all the async calls accept an additional callback, but it was a very small price to pay for extended testability. An example asynchronous test would look like this:

```

describe("Asynch Mocha Test", function () {
    it("look like this", function (done) {
        myAsyncCall(param1, param2, function () {
            // This callback expected to run when async returns

            // Test ...
            // Test ...
            // ...
            done();
        });
    });
});

```

Overall Mocha offered more flexibility and better asynchronous support at the cost of additional setup. This increased flexibility came in the way of being able to select our own assertion library so that our tests looked the way we wanted them. The cost of that was the need to find our own assertion library among the dozen that are out there. We prefer Mocha's method of asynchronous testing because it required considerably less typing on our part and was more easy to follow and read than a chain of method calls.

4.1.2.1 Assertion Library of Choice: Chai.js [7] For the assertion library we chose Chai.js. It offers the most flexibility for writing assertions that we'd found. It provides access to an Expect API, a Should API, and a more general Assert API. The Expect API provides assertions in the form of `expect(myVar).to.equal(10);` or `expect(myArr).to.have.length(10);`. The Should API has assertions of the form `myVar.should.be.a('number');` or `myArr.should.have.length(10);`. Finally, the Assert API has the more tried and true `assert.equal(myVar, 10);` and `assert.lengthOf(myArr, 10);`. The various APIs combined with Mocha's flexibility led to a nice developer experience that can appeal to a variety of different developers.

4.1.3 QUnit [22]

QUnit is an Assertion style framework developed by the jQuery Foundation, the makers of jQuery, an immensely popular framework for DOM manipulation and access. QUnit is a bit more developer-centered and not as "feel-good" as the other frameworks but has, so far, been the most easy to use and has the most thorough test reporting of any framework out of the box. Mocha and Jasmine show how many "specs", the `it` functions, have passed, which is great information, but QUnit takes it a step further and breaks down each spec into the asserts contained within:

Assertions:

Module	Test	Assertion	Result
Task Spec	All tasks can be shown	Correct page is being rendered.	ok
		Both tasks are being shown.	ok
	A task can be added	The additional task is present.	ok
		The exact task we added is there.	ok
	All tasks can be marked completed	Both tasks are marked complete.	ok
		marked one as completed	ok
	All completed tasks can be shown.	The one we marked completed is shown.	ok
		The correct page is shown.	ok
	A single task can be marked as comple_	Our one task is marked completed.	ok
	A task can be deleted.	Only one task remains.	ok
		Page was sent the delete code, 200	ok

Figure 5: Each individual spec is further broken down into the assertions within it.

An interesting feature that QUnit has that we'd not seen elsewhere is the ability to specify how many assertions are expected to run in the test. This is incredibly useful for callbacks. By setting an expected number of assertions, tests can no longer fail silently just because an assertion was never run. If the QUnit test finishes and the assertion count is not what was specified, it will fail the test and let the developer know fewer assertions were called than expected.

QUnit isn't perfect however, QUnit tests are not grouped into suites by default and the syntax for grouping them leaves much to be desired. In order to group tests into "modules", as QUnit calls them, the developer writes code like the following:

```
QUnit.module("Module A"); // All code following is grouped into Module A
test("Addition", function () {
    equal(2 + 2, 4, "Two + Two = Four");
    equal(3 + 5, 8, "Three + Five = Eight");
});
test("Subtraction", function () {
    equal(4 - 2, 2, "Four - Two = Two");
    equal(2 - 2, 0, "Two - Two = Zero");
});

// More tests....

QUnit.module("Module B");
// Tests under Module B

// ... Etc
```

It is difficult to visually see where one module starts and ends and so grouping and finding the right place to add a test can be a chore.

QUnit's handling of asynchronous tests is similar to Mocha's except instead of adding a parameter to the `it` function, the developer calls the `asyncTest` function instead. Finally, the developer calls `start()` when the asynchronous test is finished. Example:

```
asyncTest("Async Test", function () {
    myAsyncCall(param1, param2, function () {
        // This callback expected to run when async returns

        // Test ...
        // Test ...
        // ...
        start();
    });
});
```

QUnit was a particularly strong contender when we saw the table that was produced after running the tests but its shortcomings led us to keep looking elsewhere. QUnit's handling of asynchronous tests is nice because they are more explicitly stated for the casual read-through instead of an extra parameter or set of functions within the test itself.

4.1.4 Intern[47]

Intern is a hybrid unit testing and browser testing framework that aims to be an all-in-one testing framework. Intern has unit testing, browser testing, and code coverage information all built into the same tool. Intern runs on its own server which give it access to the internals of the project, but caused problems when attempting to interact with a separate server running the Node.js application. We were unable to get the browser testing portion of Intern to work because of the inability to set up cross-server communication easily and properly.

Intern uses Asynchronous Module Definition (AMD) form which is something we'd not seen before. Where this comes into play is merely in what used to be `require(...)`. Instead of separately requiring modules, the entire program is wrapped in a `define([...])` block that passes modules as parameters to an overarching function. That will make more sense after an example:

```
// Intern
define([
  'intern!object',
  'intern/chai!assert',
  'app/hello'
], function (registerSuite, assert, hello) {
  // Test code here
});

// CommonJS, what Node.js uses
var registerSuite = require('intern/object');
var assert = require('intern/chai/assert');
var hello = require('app/hello');

// Test code here

// Another way to write CommonJS so that it's not global scope, used frequently
(function () {
  var registerSuite = require('intern/object');
  var assert = require('intern/chai/assert');
  var hello = require('app/hello');

  // Test code here
} ());
```

Using either style you get access to the modules outlined. AMD form offers a tighter, more straightforward way to look at module loading in addition to being better supported and having fewer gotchas. [38] Intern offers support for CommonJS modules but strongly recommends switching to AMD styles for your project. We chose to keep out CommonJS

format because of the maturity of the evaluation project.

Intern uses Chai.js for assertion methods but has it built-in as opposed to Mocha in which the developer had to get it themselves. This gives the developer access to Expect, Assert, and Should testing methods. In addition to the various assertion styles, Intern offers three different testing interfaces: TDD, BDD, and Object. These testing interfaces effect how the general layout of test code looks. Here's a very quick example of the different layouts:

```
// TDD
suite('TDD Suite', function () {
  test('test 1', function () {
    ...
  });
})

// BDD
describe('BDD Suite', function () {
  it('should be test 1', function () {
    ...
  });
})

// Object
registerSuite({
  name: 'Object Suite',

  'test 1' : function () {
    ...
  }
})
```

The ability to mix assertion styles with testing interfaces provides the most diverse ways of writing tests that we've found. Some testing interfaces seems to fit better with certain assertion styles, but there's no restrictions on which can be used with which.

Intern uses the WebDriver API to communicate with Selenium to drive the automated browser testing portion of Intern. Selenium and the WebDriver API will be discussed in more detail in the next section "Browser and UI Testing".

Intern's unit-test versus browser test seperation is termed, by Intern, nonfunctional versus functional testing. Nonfunctional tests are those that run without a browser while functional tests are WebDriver tests.

Intern includes configurable code coverage metrics on tests by default. Running unit tests gives the number of lines covered (ran) by the tests. Code coverage is provided by another Javascript library, Istanbul, created by Krishnan Anantheswaran [4].

Intern offered a one-stop shop for all of our testing needs but fell short when it came

to configuring it for our Node.js server setup. If we had a more traditional web application setup with a separate server that served normal HTML pages, Intern would have likely been our go-to tool for all of our testing needs.

4.2 Browser and UI Testing

Web applications, for the most part, require user input in order to be useful. The way the user interacts is through the outward facing web pages. Testing of a web application would be incomplete without testing the interactions and elements of these web pages. It's important that the inner workings are running the right code and producing the right answer, but it's almost just as important that the parameters get passed correctly and from the right places on the web pages. Also, it is important that navigation work as expected or a user will become frustrated and lose interest in the application.

What's important in browser testing is repeatability and being able to test on multiple browsers. Without the ability to automatically repeat tests, developers are left manually clicking through web pages every time anything changes. In addition, because there are so many different browsers and browser versions in the wild, it's important that the web pages run correctly (and predictably) on as many of them as possible.

4.2.1 Selenium[44]

Selenium is an automated browser input runner. It is one of the only automated browser runners and certainly the most prominent. A developer can write code that will cause Selenium to drive a browser's input and assert expectations. On its own Selenium has very minimal testing usage. It lacks a test reporting mechanism and only offers output in two forms: 1) No output at all, a browser window will open, input will be generated, and the window will close, or 2) the browser will exit with a cryptic error printed to the console. It is up to the developer to include print statements as to what is happening at any given time, what test is being run, and what the results are.

Generally, all browser testing tools will be using Selenium behind the scenes and doing the housekeeping automatically. This is the case for both of the web testing tools that will follow this description.

Selenium has support for running a number of different browsers.

Having the ability to automatically run browser tests at any point takes a huge burden off of the web developer and allows for UI testing in addition to behind the scenes functionality testing. Whenever code is changed that may affect the UI flow or functionality, Selenium tests can be run without the need for developers to manually click all of the HTML elements and input information in the relevant places.

Selenium's node incarnation is out-of-date and as such we had to change the selenium package to use the latest version of Selenium that works with the latest version of browsers. This required us to edit some of the javascript files that ran Selenium from the console

using node. At the time of this writing, the Selenium node package was using Selenium 2.20 from over two years ago. The latest Selenium version is 2.40.

Selenium was written with maximum functionality and language flexibility in mind and therefore can be incredibly verbose.

The Selenium WebDriver has APIs for a variety of programming languages including: Javascript, C, Java, and Python.

4.2.2 Nightwatchjs[40]

Nightwatch is a Javascript testing library that wraps Selenium in an effort to make browser testing more concise and straightforward. Nightwatch uses Node.js' assertion library paired with Selenium's WebDriver to both run browser automation and also report test results.

Tests are written as a long chain of function calls that produce an easy-to-read series of browser events. For example, the following code segment opens a browser, goes to google.com, makes some assertions, writes "nightwatch" in the search box, clicks the search button, and then makes a few more assertions:

```
module.exports = {
  "Demo test Google" : function (client) {
    client
      .url("http://www.google.com")
      .waitForElementVisible("body", 1000)
      .assert.title("Google")
      .assert.visible("input[type=text]")
      .setValue("input[type=text]", "nightwatch")
      .waitForElementVisible("button[name=btnG]", 1000)
      .click("button[name=btnG]")
      .pause(1000)
      .assert.containsText("#main", "The Night Watch")
      .end();
  }
};
```

As one can see, just reading the code line by line paints a picture of exactly what is happening and what is being checked for.

The same code in normal Selenium is quite a bit more verbose:

```
var assert = require('chai').assert,
    test = require('selenium-webdriver/testing'),
    webdriver = require('selenium-webdriver');

test.describe('Google Search', function() {
  test.it('should work', function() {
    var driver = new webdriver.Builder().
```

```

    withCapabilities(webdriver.Capabilities.chrome()).
    build();

    driver.get('http://www.google.com');
    driver.findElement(webdriver.By.name('q')).sendKeys('nightwatch');
    driver.findElement(webdriver.By.name('btnG')).click();

    driver.sleep(1000);
    driver.findElement(webdriver.By.id('main')).getText().then(function (text) {
        assert.isTrue(text.indexOf("The Night Watch") !== -1);
        driver.quit();
    });
});
});

```

In addition to being more verbose, I've had to add an assertion library and this code has to be run using Node.js and the Mocha module instead of using the Selenium command.

Nightwatch is Node.js module, and as such can easily installed and run via npm and the command-line.

Though a small feature, Nightwatch has the ability to either assert or verify tests. The difference is that assertions will stop execution while verifications will merely log the failure and continue on. Verification can be useful if the developer wants to get an update on how many things are currently failing as opposed to having the entire test stop executing immediately after a single failure.

Nightwatch, as of this writing, has the ability to open Firefox, Chrome, and Internet Explorer. As of February 2014, Chrome is used by 46.69% of internet users, Internet Explorer is used by 24.39% and Firefox is used by 20.78%.[48] Together that accounts for 91.86% of all internet users.

4.2.3 Headless Browsers

In an attempt to speed up UI testing, which is notoriously slow, there have been efforts to create “headless” browsers which behave like normal browsers except that there's no display element. They emulate the DOM and skip the display phase. We opted to stick with visual browsers due to time constraints and the enjoyment of watching the tests execute.

Nonetheless, we still looked at two options that we will outline briefly below for other developer's reference.

4.2.3.1 PhantomJS[16] The reigning king of headless browsers is PhantomJS. PhantomJS is a headless browser written using WebKit, which is what most of the web is based on. This means that running PhantomJS is nearly indistinguishable from a normal browser to any testing code.

How important is having these? Mention as an option but don't go into detail?

There are, however, a couple problems with PhantomJS. It is generally fairly verbose when it comes to writing the scripts that run your commands. It is also not a testing tool, it is a headless WebKit. The distinction here is that PhantomJS does not have any assert functionality built-in and thus requires the developer to find and use their own assertion library.

The difference between PhantomJS and ZombieJS, which I talk about in the next section, is akin to the difference between Selenium and Nightwatch. PhantomJS is heavier duty and requires an additional set of testing tools whereas ZombieJS is lightweight and has assertions built-in.

4.2.3.2 ZombieJS[5] ZombieJS is often the other go-to for headless browsers. ZombieJS is not a WebKit implementation but rather a simple DOM emulator. This means that certain tests may not be as accurate as they could be with PhantomJS.

Where ZombieJS shines is in its simplicity and brevity. Much like Nightwatch, calls are chained together to create a flow of calls that generally end with assertions.

```
// Zombie.js website's example code
var Browser = require("zombie");
var assert = require("assert");

// Load the page from localhost
browser = new Browser()
browser.visit("http://localhost:3000/", function () {

    // Fill email, password and submit form
    browser.
        fill("email", "zombie@underworld.dead").
        fill("password", "eat-the-living").
        pressButton("Sign Me Up!", function() {

            // Form submitted, new page loaded.
            assert.ok(browser.success);
            assert.equal(browser.text("title"), "Welcome To Brains Depot");

        })
    });
```

Developers like ZombieJS syntax so much that there exists a wrapper that uses ZombieJS's API around PhantomJS. [49]

ZombieJS offers a lightweight solution to headless browser testing, but does not have the same depth of features as PhantomJS or the same assurances that your tests do exactly what they would do in a browser with a head.

4.3 Continuous Integration Frameworks

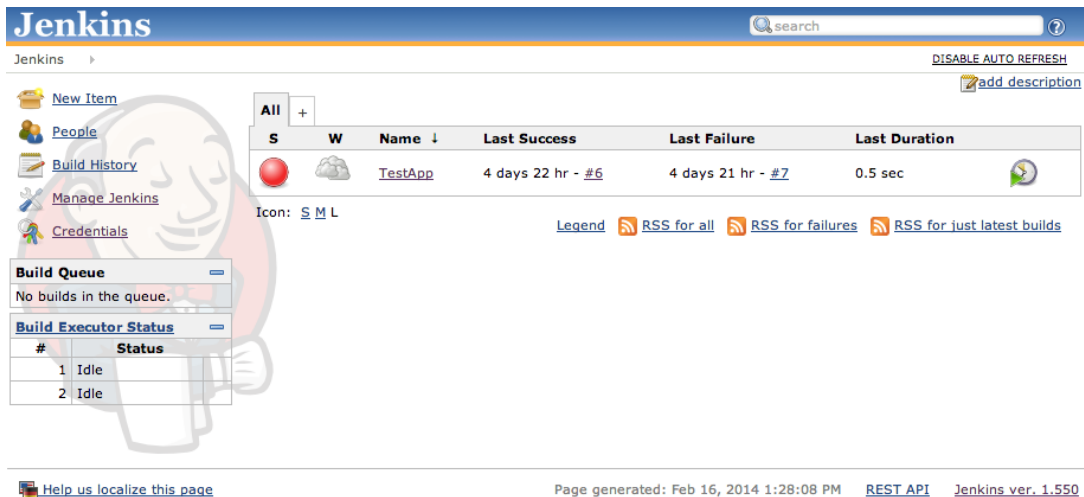
Continuous Integration frameworks attempt to make development easier by running specified actions, usually tests, after a change to the project. This change can be simply saving a modified file, or, more usually, an action like pushing a change to a repository.

Continuous Integration frameworks, therefore, need to be flexible in the type of project they support, customizable in their chosen environments to support differing projects, and have a simple way to connect to repositories or file systems.

4.3.1 Jenkins [23]

Jenkins is the first of the continuous integration frameworks that we tried out. It was created by developers of a previous tool called Hudson when a dispute with Oracle led them to fork the Hudson project and rename it Jenkins. As of February 2014, the Jenkins organization has over 1,100 public repositories and 570 members. [20] Jenkins has, according to their website, two main focuses: (1) building and testing software continuously and (2) monitoring executions of externally-run jobs. The first focus is the general goal for any CI framework. The second focus is of note and will be discussed after the initial evaluation.

Jenkins is a series of web pages that are set up to run on a given server. This is what it looks like:



The screenshot shows the Jenkins web interface. At the top is a blue header with the 'Jenkins' logo and a search bar. Below the header is a sidebar with links: 'New Item', 'People', 'Build History', 'Manage Jenkins', and 'Credentials'. The main content area displays a table of builds for a job named 'TestApp'. The table has columns for 'S' (status), 'W' (workspace icon), 'Name', 'Last Success', 'Last Failure', and 'Last Duration'. A single build is listed with a red status icon, a cloud workspace icon, and a duration of '0.5 sec'. Below the table are links for 'Icon: S M L', 'Legend', and three RSS feeds: 'RSS for all', 'RSS for failures', and 'RSS for just latest builds'. On the left, there are two expandable sections: 'Build Queue' (showing 'No builds in the queue.') and 'Build Executor Status' (showing two executors in 'Idle' status). The footer contains a link to 'Help us localize this page', the text 'Page generated: Feb 16, 2014 1:28:08 PM', a link to the 'REST API', and the version 'Jenkins ver. 1.550'.

S	W	Name ↓	Last Success	Last Failure	Last Duration
		TestApp	4 days 22 hr - #6	4 days 21 hr - #7	0.5 sec

Most companies give Jenkins a sub-domain on their web pages.

Jenkins comes bare-boned but is highly customizable and has a wide array of plugins. The plugins available integrate Jenkins with various services, integrate Jenkins with source control systems, and even just change the look of the Jenkins page or how it behaves. As an example, Jenkins can integrate with GitHub to run builds and tests when changes are pushed to a branch.

Because Jenkins comes bare and runs on the system of the developer's choosing, it can run any project that is set up. The downside is that there is no support out-of-the-box unless the system already has the necessary compilers, interpreters, applications, and/or libraries installed.

One of the exceptional strengths of Jenkins is that the developer is only limited by the hardware of the machine running Jenkins. Any increase in computing power will be directly usable by Jenkins. Jenkins can run as many jobs as the server can handle at one time.

The other use for a tool like Jenkins is the ability to start services like cron jobs or similarly long-running tools and status monitors and then collect the output for them and allow the developer to check on them at any time.

Jenkins is an exceptionally strong tool for CI, but requires immense overhead in setting up the environment to run a project in. However, once the server has the necessary applications and libraries, the builds and tests will run without issue. Jenkins can run an "unlimited" number of builds, tests, and processes with the only limiting factor the hardware that it is running on. It is also 100% free and can connect with public and private repositories alike.

4.3.2 Travis [8]

Travis is an online continuous integration framework with 29 members and 64 repositories at the time of this writing. Unlike Jenkins, Travis is all online and includes a dashboard from which the developer can monitor builds and processes. Travis has a free version which is fully-featured but requires the project to be open-source and a paid version which gives the developer more processing power and access to private repositories.

Travis comes with internal support for eleven different languages including NodeJS, Ruby, Java, and Python. Travis has incredibly thorough documentation and a wealth of built-in services to work with projects if needed. These services include databases, notifications, deployment options, and GUI/Headless browsers. Travis supports all of the most popular databases: MySQL, PostgreSQL, MongoDB, CouchDB, and a myriad of others. It offers a number of notification options ranging from Email to IRC to HipChat (a chat service from Atlassian). Travis also offers deployment to just about every sort of provider there is, from Heroku to AWS S3 to RubyGems and also provides the ability to create custom web hooks to deploy to a provider not supported natively.

If at any point the project requires something that Travis does not support natively, Travis gives the permissions necessary to install any sort of libraries or services that can be installed with apt-get on an Ubuntu machine.

Travis requires that projects be hosted on GitHub, but by doing so the setup is incredibly easy. The developer signs into Travis using GitHub credentials and Travis pulls up all of the developer's repositories automatically.

Getting a project running with Travis is fairly easy; just create a yml file in the reposi-

tory with certain lines pertaining to the language, needed services, commands to run, etc. We were able to set up everything we needed in six simple lines.

In order to get access to private repositories, Travis offers several pricing plans, the cheapest of which is \$129 [9] for their “Startup” plan. That’s a bit steep, but it does offer unlimited private repositories, unlimited collaborators, and two concurrent jobs.

From what we gathered, Travis appears to be targeted at companies, not individual developers. Travis requires developers to host their projects on GitHub and only supports one process running at a time. However, Travis is incredibly easy to set up, does not require it’s own server, comes with impressive built-in features, and works wonderfully with open-source repositories on GitHub.

4.3.3 Semaphore [36]

Semaphore is another online CI framework that connects to GitHub to bring in projects. Semaphore is made by RenderedText, a software consulting team. Semaphore connects with public and private repositories but requires a subscription after 30 days no matter which repositories are used.

Semaphore was originally designed for only Ruby applications but in January of 2014 they announced official support for Node.js and Javascript and as of February 26th they now offer official support for Clojure projects.[37] These types of projects require no setup or files in the repository, Semaphore automatically detects project type and auto-fills build steps. Their “supported stack” page also lists C/C++, Java, Perl, and Python, but those must be configured manually.

Semaphore supports most databases and has a number of testing frameworks built-in. Once again any missing libraries or dependencies can be installed using apt-get. It would appear that deployment support is only available for Ruby projects.

Semaphore’s pricing scheme is different than Travis’ as it focuses on an individual, not a company. Their lowest cost subscription is one project (public or private) for \$14 per month. They offer more subscriptions with the ability to have more projects and concurrent processes in higher tiers.

Semaphore is still ruby-centric and requires a subscription. As our goal is to find a framework that is free for open-source projects, Semaphore was out automatically. It is included here because it’s well-known and fairly established, especially in the ruby community.

4.3.4 MagnumCI [25]

MagnumCI is another online CI, it appears that Jenkins is one of the only “off-line” CI systems. MagnumCI has every feature available for free, however, it’s still in public beta and although prices haven’t been released, the team making MagnumCI has already stated

that there will be price tiers. It is not unreasonable, though, to believe that MagnumCI will be free for open-source projects, and that's what we care about.

MagnumCI offers support for Ruby, PHP, Node.JS, and Go. MagnumCI, much like Semaphore, will attempt to auto-detect commands that it should run based on the project structure and language. All settings are overridable in a yml file.

MagnumCI has all the popular databases built-in, many of which run automatically on startup. MagnumCI also has two headless browser frameworks built in for running any browser testing that needs to be done. The developer can, again, use apt-get to install missing packages. MagnumCI does not, at current, provide automatic deployment options given a successful build.

Where MagnumCI shines is in support of different repositories. MagnumCI offers support for pulling from GitHub, BitBucket, GitLab, and a few others. It does require a bit more setup to get repositories connected, but what it comes down to is a couple lines of copied and pasted text within the repositories' settings.

Right now MagnumCI offers impressive features, private repositories, and extensive repository sources, all for free. It is hard to say that MagnumCI is the best because those features may end up costing money soon after this paper is completed. It is included in this review because it seems to be an up and comer that could prove quite strong.

4.3.5 Drone.io [41]

Drone.io is an impressive online CI framework that offers substantial language support, is free for open-source projects, and offers the most reasonable pricing scheme for private repositories.

Drone.io supports eleven languages including: C++, Dart, Go, Haskell, Java, Node.JS, Ruby, and others. Setup requires the developer to specify language, as opposed to the automatic tools provided by the previous two CI frameworks, but will auto-generate build commands.

Drone.io offers support for most of the databases we've heard of, has a variety of deployment options, and is entirely configurable from the online dashboard. It also offers the most browsers and headless browsers for use with web page testing tools such as Selenium and PhantomJS.

Drone.io is able to connect to GitHub, BitBucket, and Google Code repositories to get projects automatically. Setup is simple, select the source of the repository (GitHub, BitBucket, or Google Code), select the repository from the list it gives once it connects, select the language of the repository, edit the default build script, and the project is ready to go.

Drone.io offers unlimited free open-source repositories. The next step up offers 5 private repositories for \$25 per month and for \$50 per month, unlimited private repositories. They also offer higher tier plans that allow for concurrently running builds.

Drone.io is a strong contender because of its language support, deployment options,

and ease of setup. Where it lacks is in its notification options. The only notification option at present is email while other CI frameworks offer a number of other notification options for various messaging services.

4.4 Misc Frameworks

In addition to the Unit Testing, Continuous Integration, and Browser Testing frameworks outlined above, we came across a number of incredibly helpful frameworks that didn't quite fit into any of the above categories.

4.4.1 Sauce Labs[43]

Sauce Labs is an online service that gives you access to a variety of browsers and browser options. They outline the options available to the developer as follows: Selenium, Javascript, Mobile, Manual. Sauce Labs will start up a Selenium server and run the developer's Selenium tests, they can run unit tests using their Javascript service, they allow access to a variety of mobile device browsers and resolution, and they offer the ability to manually look at a web application in different browser configurations.

Every CI framework we looked at had built-in configurations for Sauce Labs so that one of the CI steps could be running your tests on Sauce Labs.

Sauce Labs is an incredibly helpful tool for making sure a web application behaves as expected on a wide spectrum of devices and configurations and saves the developer the hassle of acquiring multiple devices, operating systems, and browsers.

4.4.2 Istanbul[4]

As already mentioned briefly, Istanbul is a tool used for code coverage metrics. It was built into Intern.io but getting it running with other frameworks is easy as well.

Getting Istanbul to output code coverage is as easy as installing it and then calling the `istanbul` function from the commandline with your test executable. For example: `istanbul cover mocha -- . -R spec` will run the mocha tests in the current directory and output coverage statistics. To break that down a little further:

`istanbul cover` tells istanbul to run coverage on the tests.

`mocha --` specifies that the executable to be run is mocha and that the developer is done giving commandline args to istanbul, that's the `--`

Finally, `. -R spec` is just the flags for the mocha command to run the appropriate test suites.

Istanbul offers an incredibly easy way to get code-coverage on your tests and is painless to get working right away.

4.4.3 Grunt.js[6]

Grunt.js is a task-runner for Javascript. What Grunt does is automate routine tasks like linting, testing, copying files, and updating modules. The developer writes Grunt scripts that outline the necessary tasks and then can run them on the commandline with the `grunt` command.

Grunt is widely supported and has all kinds of built-in or installable plugins that make writing common tasks nearly trivial. As an example, a few built-in grunt tasks that we use include: `jshint` for linting, `uglify` for code minimization, and `watch` for automatically re-running tests when certain files change. Each of those plugins requires very little configuration and they're ready to be used.

Of particular interest is that last feature, `watch`. We are particularly fond of grunt's `watch` capability because it means that we have tests running near constantly that give immediate feedback to the developer when they change existing code or write new code after having written a test for it. As long as all the tests are passing, the developer has a good idea that the changes made are regression-tested as well as freshly tested and are ready for publishing or integration.

4.5 Misc Tools

4.5.1 Node Inspector[10]

Node Inspector is a GitHub project for debugging Node.js server-side code. The inspector runs in the browser and is started by running `node-debug` instead of `node` when starting the application. Simple as that. Then you can set breakpoints using the normal browser developer tools.

Currently Node Inspector only runs on Chrome and Opera but since those browsers are available on all OSes, it shouldn't be a problem for developers to use.

5 Implementation

5.1 Chosen Toolset

After evaluating a number of different combinations and tools, we came to the following setup:

- Unit Testing: Mocha with Chai.js
- Continuous Integration: MagnumCI
- Browser Testing: Nightwatch
- Misc:

- Coverage: Istanbul
- Task Runner: Grunt
- Package Manager: NPM (sort of a given when working with Node.js)

5.1.1 Mocha with Chai.js

Our choice of using Mocha with Chai.js was made because Mocha offers the flexibility of choosing one's own assertion library while still being structured in a way that all developers can read the code.

Mocha was chosen over Jasmine because of Mocha's additional flexibility and the way that Mocha handles asynchronous tests.

Mocha was chosen over QUnit because of the ability to structurally portray test suites. Although we loved QUnit's output, the ability to scan code more efficiently won out in the end.

The final decision came down to Mocha versus Intern. Mocha was eventually chosen over Intern because of the ability to choose an assertion library and because of the issues we had with Intern. Both frameworks offer impressive flexibility regarding the way in which tests can be structured and written. Mocha ended up allowing for a bit more flexibility and caused no problems for us. Perhaps with more time and patience we would have chosen Intern, but Mocha served our needs well.

5.1.2 MagnumCI

With all of the features that MagnumCI currently offers for free, we couldn't pass it up. We are still confident that the features needed by an open-source developer will continue to be free after MagnumCI ends its open-beta period.

Quite frankly, we didn't have the time or interest in setting up our own Jenkins server. Given enough resources and time, Jenkins will surely beat all of the rest of the CI frameworks, however, there's something to be said for ease of setup and total time to first interaction with the tool.

MagnumCI was chosen over Travis because our evaluation application was not on GitHub. Travis is still an incredible CI framework that has some of the best documentation we've seen.

Semaphore was out of the running due to its lack of support for JavaScript and Node.js projects as well as the GitHub requirement.

Drone.io, while offering a number of repository location options, did not offer the ability for a repository to be hosted at a custom location.

Overall, the CI frameworks were all impressive and usable. The main driving factor was simply that we didn't have our application on any of the more common repositories and at the time of this writing our application was closed-source. We have every intention of opening it to the public as soon as we get a version that is polished.

5.1.3 Nightwatch

The choice for Nightwatch was an easy one. First and foremost there were not very many options and the options we did have sorted themselves out.

Nightwatch, being a wrapper around Selenium already, was chosen over Selenium because of its concise way to express tests. We have yet to come across any situation that is not covered with Nightwatch as the developer of Nightwatch has seemed to implement nearly all WebDriver calls, even if some are undocumented.

Nightwatch was chosen over Intern's UI testing tool because of the issue with Intern running its own server. If Intern did not impose its own server for what it calls "functional" testing, we would probably have gone with Intern for both Unit and UI testing but it did not make sense for us to split up Intern's functionality since we couldn't use it all.

5.1.4 Miscellaneous Tools

Generally the miscellaneous tools were chosen because they were already integrated into our project or were the obvious choices for other reasons.

Istanbul was chosen because we had seen it in use with Intern and after minimal research found that it was easy to set up and get running.

Grunt was already a part of the evaluation application prior to the testing effort we set out on. With a bit of looking into Grunt, it was an obvious choice to keep it as our new functionality was easily integrated into the existing grunt scripts.

NPM is the go-to choice for acquiring modules for Node.js applications.

6 Product

6.1 Our Setup

The following are the specs of the laptop that we used for the entirety of this research and implementation:

- Model: 15-inch Macbook Pro, Late 2008
- OS: OS X 10.9.2, Mavericks
- Processor: 2.53 GHz Intel Core 2 Duo
- Memory: 4 GB 1067 MHz DDR3
- Graphics: NVIDIA GeForce 9400M 256 MB

The point of listing the laptop specs is to show that our research and workflow can be used on a very run-of-the-mill laptop with generally sub-standard specs. Or at the very least, an average, everyday laptop.

To carry out the development we used a combination of the OS X terminal, Sublime Text, and Netbeans. The terminal was used mainly for starting Node services, testing tools, Grunt scripts, and updating modules. Sublime Text was used to write the unit tests and the UI tests. Netbeans was used to look at the overall project as it had already been set up in Netbeans and was easy to browse therein.

6.2 Tool Setup

Getting the tools up and running is a fairly simple task. It involves installing modules with NPM, writing or copying a few configuration scripts, and then knowing the right commands.

We decided to install everything globally because it made writing our scripts easier, but it's not a requirement. Everything can be done with locally installed node modules, instead of using the command by itself the developer would just need to supply the full path to the executable, which is generally `PATH\TO\REPO\node_modules\MODULE_NAME\bin\MODULE_NAME` or something similar.

6.2.1 Setting up Mocha

To install Mocha, type into the terminal: `npm install -g mocha`. The `-g` means to install the module globally.

Next, to install Chai, type into the terminal: `npm install -g chai`.

To get Chai running with Mocha, just `require(...)` the appropriate chai assertion method in the test and start writing. For example, to get the Assert library: `var assert = require('chai').assert;`

After the tests have been written, the developer can run them using the `mocha` command. There are a number of command-line arguments available to the `mocha` command which are well documented and which we won't talk about here. We will talk about the one that we used consistently, that being `-R`. The `-R` arguments allows the developer to specify how the test results will be formatted. Mocha offers seventeen different reporters with more reporters available online and the ability to create custom reporters.

6.2.2 Setting up MagnumCI

Setting up MagnumCI is painless. Visit “magnum-ci.com”, create an account, sign in, click on “Add a New Project”.

Here is the new project page:

MagnumCI

[Dashboard](#) [Builds](#) [Doc](#)

Create a New Project

Name

Ex: simple-sinatra-app

Repository URL

Ex: git@github.com/username/repo.git

Integration

Select one ▼

Source control

Git ▼

Project type

Generic ▼

Create Project

[Cancel](#)

Fill out the fields with the appropriate values for a given project and click “Create Project”. Depending on what the setup is for your project, you’ll see slightly varying web pages after you create your project. One thing in common, though, is the button at the bottom of the steps list that says “Customize Build”. After following the instructions, click “Customize Build”:

Basic Setup

Build Configuration

Collaborators

Web Hooks

Addons

Deployments

Build Configuration


Build on branches


Any Branch


Runtime Version


Select version


Resources


 PostgreSQL


 MySQL


 MongoDB














Environment variables

One per line. Example: FOO=BAR

Build Steps

Before installation:

Dependencies installation:

Before test execution:

Test suite commands:

After test suite execution:

Save Changes

[Back to Project](#)

There are a number of configuration options and panes that can be accessed on the left side. Our focus was making sure the “Build Configuration” tab was properly set up. This involved selecting the correct branch, the correct Node.js runtime version, the correct database from the “Resources” area, and adding the appropriate commands in the build steps. In our case we wanted “mongoDB” and “Xvfb” for UI testing, and then `npm install -g` in the “Dependencies Installation” to make sure our modules are available when we need them. Finally, we run Nightwatch using `nightwatch -c nightwatch.json` (the nightwatch flags will be outlined in the next section) in the “Test suite commands” box. All of that makes our build configuration page look like this:

Build Configuration

Build on branches

List branches

Runtime Version

Resources



Environment variables

One per line. Example: FOO=BAR

Build Steps

Before installation:

Dependencies installation:

`npm install -g`

Before test execution:

Test suite commands:

`nightwatch -c nightwatch.json`

After test suite execution:

[Save Changes](#)

[Back to Project](#)

After that setup, our project is ready to go. Anytime we commit from then on, MagnumCI will run and notify us of successful and unsuccessful builds.

6.2.3 Setting up Nightwatch

Installing Nightwatch consists of just installing another npm module. `npm install -g nightwatch`. In addition to Nightwatch, the developer needs Selenium. The easiest way is to just download the jar straight from the Google repository here: “<http://selenium->

release.storage.googleapis.com/index.html”. Once the developer has the JAR and places it somewhere, Nightwatch is ready to be configured and run. Nightwatch configuration happens in a json file that is pointed to when the developer executes Nightwatch.

The configuration file has plenty of options to get Nightwatch to do exactly what the developer wants, some of the more important ones include the `selenium` object that holds information about where the executable is located and what ports to run on, the `src_folders` attribute to specify a source directory, and the `test_settings` object that contains configurations for default web pages, differing environments, and some desired capabilities.

To run Nightwatch, navigate to the project folder and run `nightwatch -c path_to_nightwatch_config`. This will use the `src_folders` attribute to determine where to look for the tests, start the Selenium server, and run the tests. When debugging certain functions of even test cases, tests can be isolated with the `-t` flag: `nightwatch -c nightwatch.json -t folder1/test1.js`. Nightwatch offers a number of other command-line arguments that are available on their developer guide. [40]

6.2.4 Setting up Istanbul

As mentioned in the Istanbul section, getting it running is simple. Install Istanbul with npm: `npm install -g istanbul` and then run the `istanbul` command along with any others needed. In our case, we ran `istanbul cover mocha -- . -R spec` to run Istanbul on our Mocha unit tests.

6.2.5 Setting up Grunt

Getting Grunt set up just the way the developer wants it is likely the most involved, but there are a myriad of examples and sample scripts to help them through. Installation is the same as we’ve seen before, `npm install -g grunt-cli`. After installation the developer will need to create a “Gruntfile” which is just a Javascript file called “Gruntfile.js”. Inside the Gruntfile are all of the command names and procedures that the developer wants. As the documentation outlines: “A Gruntfile is comprised of the following parts: 1) The wrapper function 2) Project and task configuration 3) Loading Grunt plugins and tasks 4) Custom tasks”.

Here is an example Gruntfile:

```
// 1) Wrapper Function
module.exports = function(grunt) {

    // 2) Project configuration.
    grunt.initConfig({
        pkg: grunt.file.readJSON('package.json'),
        uglify: {
            options: {
```

```

    banner: '/*! <%= pkg.name %> <%= grunt.template.today("yyyy-mm-dd") %>
              */\n',
  },
  build: {
    src: 'src/<%= pkg.name %>.js',
    dest: 'build/<%= pkg.name %>.min.js'
  }
}
});

// 3) Load the plugin that provides the "uglify" task.
grunt.loadNpmTasks('grunt-contrib-uglify');

// 4) Default task(s).
grunt.registerTask('default', ['uglify']);

};

```

A more complicated Gruntfile simply consists of more tasks. Taken one at a time, it's easy to pick apart a Gruntfile. Task names can refer to one or more commands, so the command `grunt testMyStuff` may run one command, or may run five commands. The task names are the ones that are used as the first parameter of the `registerTask` call.

6.3 Workflow

In addition to just the a list of tools and installation instructions, we'd like to propose a way to use those tools that is efficient and nearly autonomous. Our goal is to bring testing tools into use with as little resistance as possible. With a bit of effort at the outset, testing can become a crucial and enjoyable part of development.

What we recommend is a unit test suite that runs on save. While this may seem excessive, it provides invaluable, instant feedback for regression assurance and new functionality. The way to set this up with the above tools is have a grunt task that watches for changes in certain files and runs a script when the files change. Here's what our Grunt task looks like:

```

module.exports = function(grunt)
{
  grunt.initconfig({
    watch: {
      test: {
        options: {
          spawn: false
        },
        files: ['**/*.js'],

```

```

        tasks: ['mochaTest']
    },
    mochaTest: {
        test: {
            options: {
                reporter: 'spec',
                clearRequireCache: true
            },
            src: ['../../../../../NetBeans/PolyXpressAuthor/test/unit/functional/*.js']
        }
    }
});

var defaultTestSrc = grunt.config('mochaTest.test.src');
grunt.event.on('watch', function(action, filepath) {
    grunt.config('mochaTest.test.src', defaultTestSrc);
    if (filepath.match('test/functional')) {
        grunt.config('mochaTest.test.src', filepath);
    }
});

grunt.loadNpmTasks('grunt-contrib-watch');
grunt.loadNpmTasks('grunt-mocha-test');

// Register the mocha test task
grunt.registerTask('test', ['mochaTest', 'watch:test']);
}

```

Our gruntfile is more complex than just that, but those are the parts that make the test-on-save functionality work. As we said earlier, Grunt has a lot of plugins that are available. We use one here to make our lives easier, `grunt-mocha-test`. It is installed, once again, through npm and gives us access to the now-built-in task `mochaTest`. The run-on-save functionality comes from the `watch` tasks which has `mochaTest` in its `tasks` array.

Once that setup is done, the developer just runs `grunt --force test` once which will execute the tests and then watch for changes and continue testing on save. The `--force` flag is used so that the grunt task continues after “failure” which would occur if any tests did not pass.

For running the browser tests, we recommend using your CI tool. With some setup on your CI of choice, you can get Nightwatch running. Most CI setups have either Firefox pre-installed or Xvfb (a windowless X11 process) and all that’s required is to install firefox if needed and start the selenium server using Xvfb is needed. For the CI setup, after installing all the prerequisites, run `Xvfb :99 -ac -screen 0 1280x1024x24 &` or something

more suited to the developer’s liking. Make sure the `DISPLAY` environment variable is exported, `export DISPLAY=:99` (the numbers need to match between the `xvfb` call and the export). Then start your selenium server, `java -jar path/to/selenium.jar` and it’s all set. Finally, run Nightwatch and rejoice.

7 Evaluation

7.1 PolyXpress (PX)[12]

Our major evaluation application was an in-house application called PolyXpress (PX). As the About page says, “PolyXpress allows you to create location-based stories, build eTours, or create restaurant guides. It is the tool that will bring people to locations in order to entertain, educate, or provide amazing deals.” Users can play or create “Stories”, each of which consists of “Chapters”. Each chapter consists of various “Events”. Each “Event” generally give the user some sort of information, in the form of text, video, audio, or a picture. After the player has chosen a story, they must physically reach certain checkpoints to unlock chapters and events. Because the application requires physical movement, it works solely on devices with GPS, generally phones and tablets only. Rather than put a series of pictures showing PolyXpress here, we will put them in the appendix at the end of the paper. To get a better sense of what PolyXpress does and looks like, we recommend looking at the appendix now.

7.1.1 PolyXpress Author

Our tests covered the Authoring portion of PX.

8 Related Work

We were unable to find any works that attempted to bring existing technologies together to create a web testing framework. Instead, we came across a number of proprietary tools and research. Compared to these proprietary tools, the framework we already laid out is free and each part is supported by a community. There is no need to contact any paper authors to get software or support. We will, however, still outline them as they provided insight into what the problems in existing frameworks were, gaps in coverage, what was important in a framework, and inspiration and acknowledgment that this is a real problem that many people are trying to solve.

We also found research into common bugs, general testing techniques, and finding bugs that greatly improved our test cases and understanding.

8.1 Testing Frameworks

8.1.1 “A Framework for Automated Testing of JavaScript Web Applications” [46]

A collection of IBM Researchers and two university students came up with a framework called “Artemis” for testing web applications. What is novel about their tool is that it generates test cases automatically based on execution patterns and feedback. They have attempted to solve the problem that writing test cases by hand is both time-consuming and often difficult. Their automatic test generation lead to an average of 69% test coverage with enough tests generated.

Though it was an interesting and somewhat effective method, 69% coverage is not great and this particular framework only covers the client-side, so it was incomplete for our purposes.

8.1.2 “A Multi-Agent Software Environment for Testing Web-based Applications” [34]

Two students and a person from Lanware Limited brought AI into the web testing world with an agent-based environment for testing web applications. In order to split testing into manageable tasks for agents to carry out, they created an ontology for web testing using XML. Each agent is set up to handle one particular kind of testing with certain test data, i.e., a unit tester with data for one or two functions. This agent may then communicate with another agent who needs the results from that test, such as a test coverage agent or agent that is keeping track of test success and failure. This communication is done through a message-passing intermediary layer and a set of brokers who shuffle messages between agents.

This system is exciting and intricate but far too complex for a general web testing toolkit that just about any developer could pick up. It was important to weed techniques like this out as we wanted a fairly easy to use and understand toolkit.

8.1.3 “A 2-Layer Model for the White-Box Testing of Web Applications” [50]

A couple of gentlemen at the Center for Scientific Research and Technology in Povo, Italy came with a model for white-box testing of web applications. (White box testing is when the developer is given access to the underlying code.) This paper breaks up the testing process into two abstraction levels, the navigation model, the way in which a webpage goes from page to page, and the control flow model, the way in which information is passed and stored on a given page or between pages. These constitute the 2-layers in the title. The navigation model represents high-level test cases, like asserting that a given link redirects to the correct location, while control flow represents low-level test cases, like making sure data is persisted between pages.

This 2-layer approach is interesting and provided us with some insight, but the system was made for PHP applications and we weren't looking to implement a brand-new system and test PolyXpress with it all in one year.

8.1.4 “Invariant-Based Automatic Testing of AJAX User Interfaces” [28]

Two individuals from the Software Engineering Research Group at Delft University in the Netherlands came up with a way to automatically test AJAX UI. The core of this idea is using a crawler to infer a flow graph. This paper outlines an plugin for an automated tool, ATUSA, for creating state validators and test suites for covering paths discovered during crawling (with a separate tool, CRAWLJAX). Their tests show that with minimal manual effort, the use of ATUSA can lead to high code coverage and error discovery.

This tool showed promise as an AJAX testing tool, but we wanted something more generic. Their crawler, CRAWLJAX, however, was very interesting and while we didn't use it, the techniques it used are used in other tools we have.

8.1.5 “JSART: JavaScript Assertion-based Regression Testing” [29]

Two students from the University of British Columbia created a tool called JSART that uses run-time code instrumentation and analysis to infer invariants that can be used for regression testing. When the code runs, JSART creates a number of test cases that will be run the next time the code executes. The point being to catch any discrepancies between old code and new code. They ran their tool on nine web systems with fairly strong results. The tool was not perfect and occasionally created inappropriate tests, but generally the tests were appropriate and useful.

This is a useful concept that could be integrated into the toolkit at a later time.

8.1.6 “DOM Transactions for Testing JavaScript” [32]

Three people from Albert Ludwigs University in Germany came up with a way to deal with test fixtures that are built and torn down frequently. When testing JavaScript, many of one's tests will change the underlying web page in some way. This is often different from testing traditional software where components are separated. In order to test JavaScript, test fixtures for setting up the page and tearing down the page are exceptionally important. Because these fixtures are being made and torn down frequently, this paper discusses a way to utilize transactional memory to quickly get the web page into the right state without constant setup and tear down.

This could be integrated with one of the existing tools we will use, but our goal was to collect existing tools that did not need alteration.

8.1.7 “Testing Web Applications in Practice” [19]

Four students from the University of Seville, Spain published a paper with an overview of web testing. This paper was more of an introduction to testing, outlining the concepts of Unit Testing, Integration Testing, Regression Testing, testing the client versus testing the server, and a few others. It also describes how to test an actual web application using PHPUnit, a library similar to JUnit, for those developers familiar with Java. When writing their tests, they suggest taking more abstract actions and breaking them up into their functional parts and testing each of those. As an example, inserting a customer into the database (albeit not that abstract) is broken into tests for the function that inserts the customer and tests for making sure the customer makes it into the database.

The paper was good refresher on testing concepts and a hands-on application of those concepts. There wasn’t any new information or tools to put toward our toolkit, but it was a good reminder and outlined some testing procedures.

8.1.8 “Contract-Driven Testing of JavaScript Code” [14]

Two individuals from the Albert Ludwigs University in Germany created a tool, JSConTest, that provides a framework for making contracts for a JavaScript program, allowing for guided random testing on inputs and outputs. Contracts in JSConTest are comments above each function that are of the form “type → type”. More complicated forms can introduce parenthesis for function types. In general, JSConTest is useful for detecting type errors on input or output based on the operations performed within a function on the input and the result being output.

Although JSConTest could be useful in a toolkit, its source is no-where to be found and there were other, more robust, tools out there for our purposes.

8.1.9 “JAWS: A Javascript API for the Efficient Testing and Integration of Semantic Web Services” [31]

A researcher from Ford Motor Company published a paper regarding an API, called JAWS (Javascript, AJAX, Web Service), to facilitate testing and integration of Semantic Web Services. Semantic Web Services (SWS) are the server side of machine-to-machine interaction via the internet. This is opposed to machine-to-human interaction such as web pages that one visits personally. SWS use markup that is detailed and sophisticated which conforms to certain standards but is not human-readable.

Another interesting API and project, but we are not focusing on SWS and so including a tool or API like JAWS would be overkill and beyond trying to keep the toolkit as simple and robust as possible.

8.1.10 “WebMate: A Tool for Testing Web 2.0 Applications” [51]

Four people from Saarland University in Germany created a tool, WebMate, for automatically navigating web-pages with dynamic content and data. WebMate’s goal is to output a usage model for testing purposes. This usage model allows tests to be focused on navigation as well as functionality. This dual usage leads to the ability to test the experience as well as the functions. In addition, WebMate can be used for Cross-Browser testing by running it on multiple browsers and comparing the results. In essence, WebMate is a web application crawler seeking to automatically suss out the layout and interaction between web pages to provide a model to the tester, be it a developer or the automatic testing capabilities of WebMate.

WebMate was another interesting project that could have seen use in our toolkit if crawlers weren’t already available in other, more mainstream, tools.

8.1.11 “Continuous Testing with Ruby, Rails, and JavaScript” [35]

The well known book by Rady and Coffin talks about Continuous Testing and has lots of examples. The last two chapters of the book, six and seven, talk about setting up a Continuous Testing environment for JavaScript using Node.js and a number of other tools. It fit almost perfectly with our project and was used as the basis of our testing setup. They use JSLint for linting, Jasmine for test cases, and Watchr for running the tests continuously. Lastly it discussed writing effective tests, which is always important.

Not all of those tools made it into the eventual final toolkit, but they were all important to look at in the context of making the overarching toolkit.

8.2 Finding Bugs

One of the most important aspects of testing is actually finding the bugs in the program. A number of sources provided guidelines and suggestions for finding bugs in web applications.

8.2.1 “Finding Bugs In Dynamic Web Applications” [45]

Seven researchers from the MIT Computer Science and Artificial Intelligence Lab wrote a paper addressing testing in dynamic web applications. Common testing techniques were not made to work with dynamic content that can be generated at the drop of a hat or click of a button. Therefore, common testing techniques are inadequate for testing dynamic web applications. This paper discusses ways to find bugs in these new-fangled dynamic web applications. They use a new technique and extended a tool called Apollo. This technique generates dynamic test cases based on concrete and symbolic execution. “The basic idea is to execute an application on an initial input (e.g., an unconstrained or randomly chosen input), and then on additional inputs obtained by solving constraints derived from exercised control flow paths.”

Although they test PHP applications, the extension to JavaScript is fairly straightforward. We didn't use Apollo, but the techniques described within allowed for more focused and productive testing.

8.2.2 “Testing Web Services: A Survey” [30]

Three researchers at the Centre for Research on Evolution, Search & Testing at King's College in London wrote a survey of web frameworks for testing web services. Their focus is on Service-Oriented Computing, a paradigm in which web sites or applications provide small pieces of functionality usable by “anyone”. These services are then combined to create more fully featured applications. What they note in particular is that when using web services, the developer rarely has access to the source code and must therefore black-box test them. Corollary to that is that the developer has to implicitly trust whoever made the service. They go through and discuss a number of frameworks and come to the conclusion that it's most certainly not a solved problem for the following reasons: The frequency of testing required, Testing without disrupting the operation of the service, Determining when testing is required and which operations need to be tested. Regarding the frequency of testing, web services often change in effort to make improvements, but those improvements can change the services and break tests. Testing without disrupting service is difficult because these services often limit accesses or simultaneous access of both a test and a running web application. Finally, determining when and what to test comes back to the trust issue. Some services provide their own testing metrics that include code coverage, number of tests, etc, but the developer must trust that the service is thoroughly tested, not just tested enough to provide interesting metrics.

8.2.3 “A Framework for Testing RESTful Web Services” [39]

Two people from the University of Dakota School of Aerospace Studies outline a framework for testing RESTful web services. RESTful web services, short for Representation State Transfer, are, in the context of web applications, a name for web APIs that conform to a set of constraints. In particular they must offer four request methods: GET, POST, PUT, and DELETE. These methods generally connect to a database or other storage system in order to store, retrieve, or delete entries therein.

8.3 Writing Effective Test Cases

Although not too distinct from *finding bugs*, writing effective test cases is its own art form. The important of writing concise, correct, productive test cases cannot be overstated. The difference between writing one test case that effective tests multiple things and writing a single test case for every little detail is the difference between developers wanting to write tests at all and skipping it due to time constraints or frustration.

Finish
reading
this pa-
per... My
bad Dr.
Haungs.

8.3.1 “Going Faster: Testing The Web Application” [17]

Two employees of Evant Software wrote an article for IEEE Software regarding the testing of web applications. Mostly this paper was an overview and implementation of Test Driven Development (TDD) and eXtreme Programming, but in the middle of the article was a good section about the difficulties of testing the web and how to deal them. A few of the remarks regarding the difficulty of web testing is that often the client-side and server-side portions of web applications are written in different languages. This wasn't an issue for our application, but is still an issue for many applications. Some of the most useful advice was: “First, test those parts of the server-side code that are not directly concerned with being part of the Web application, without involving Web peculiarities. ... Second, test those parts of the client-side code that have no server interaction. This is typically code that contains little or no business logic. ... Third, write functional tests of a low grain in the server-side language (for example, Java or C++) that simulate the request/response Web environment. ... Using this framework you can write walk-through tests that simulate the user moving through the UI”.

This step-by-step walkthrough of how to test was an excellent starting point to look at different tools and how they handled this issue.

9 Conclusion

10 Future Work

11 Appendices

11.1 ToDo App Test Code

11.2 PolyXpress Application Pictures

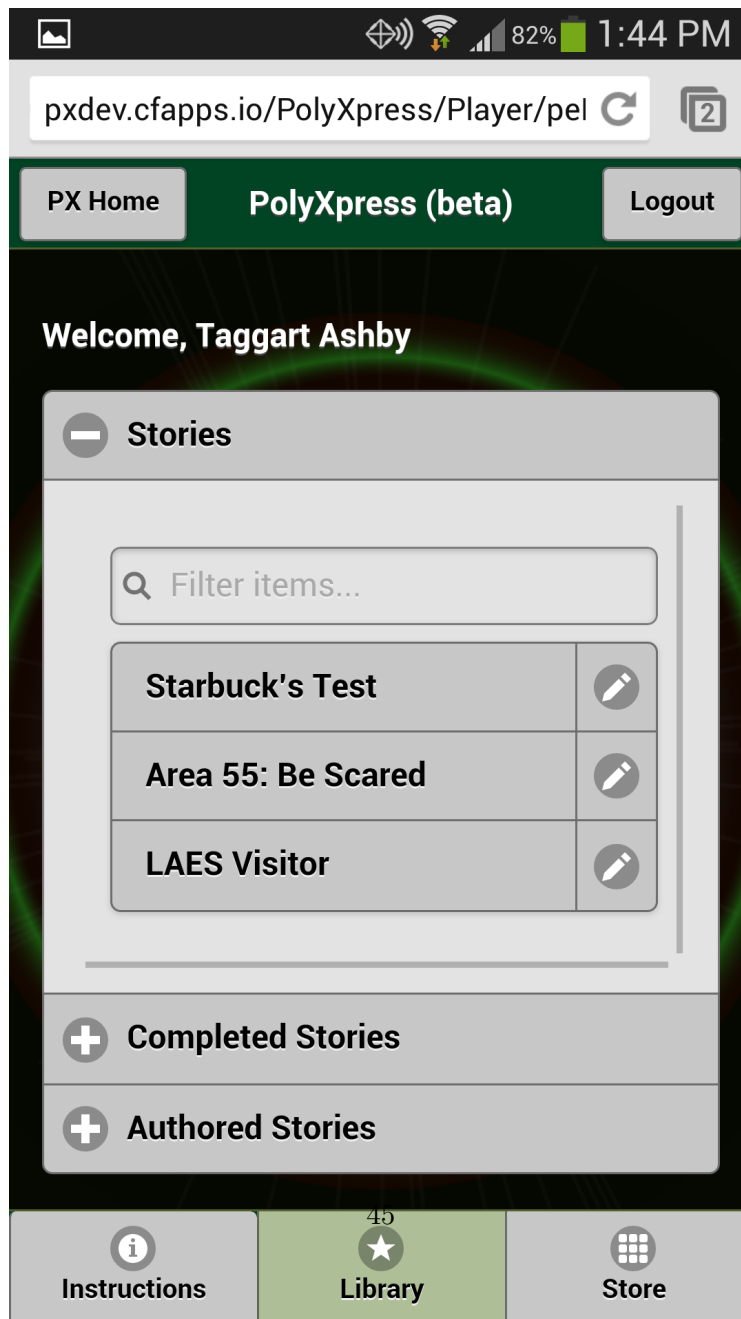


Figure 6: Story selection




Figure 7: Chapter pane with approximate distances to each chapter

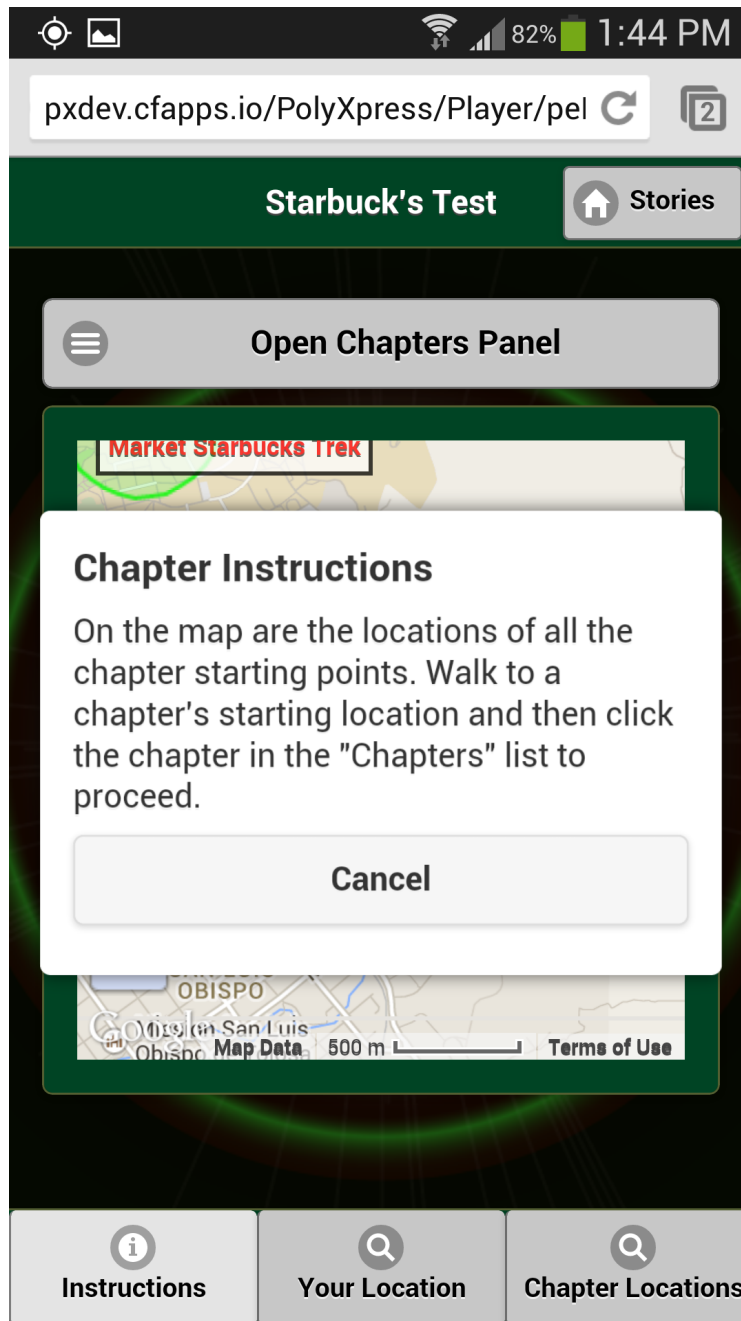


Figure 8: Each page and element of PolyXpress has instructions, here they are for Chapters




Figure 9: A map of the chapter location in relation to the player

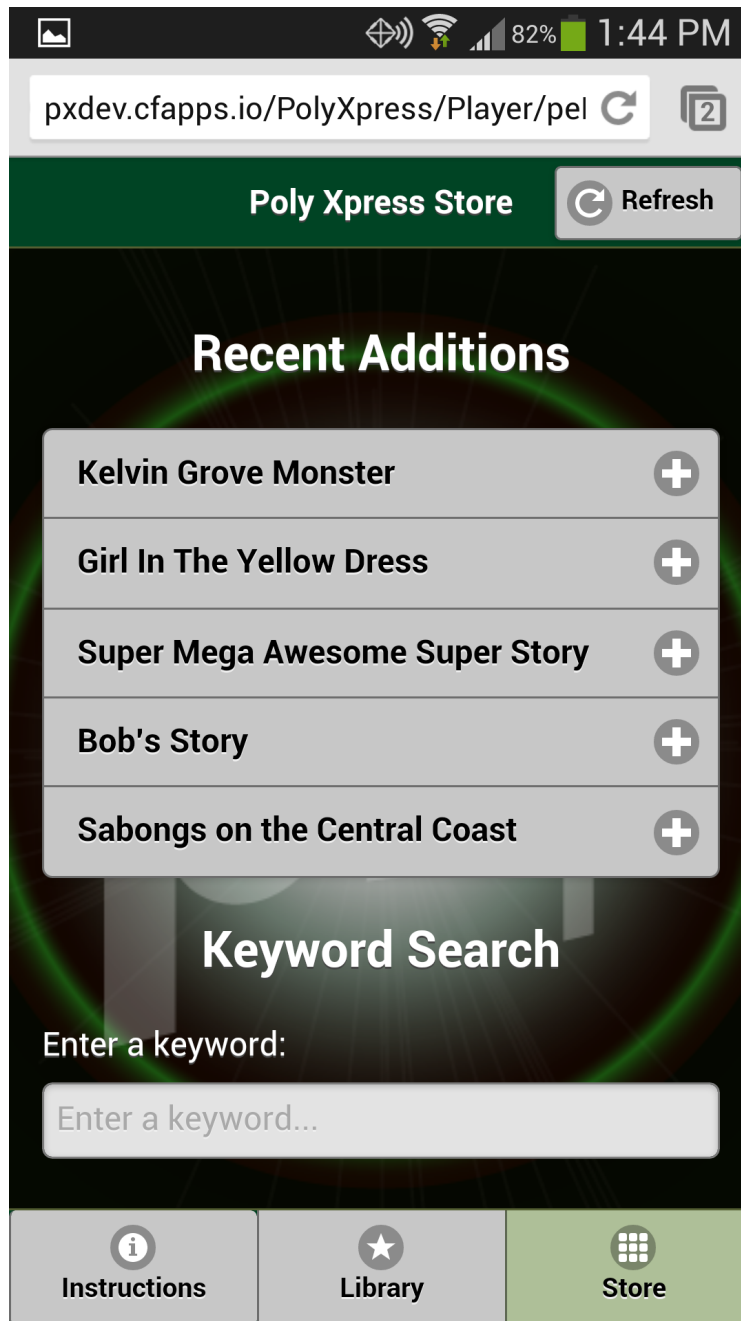


Figure 10: PolyXpress marketplace, to find other published stories




Figure 11: A map centered around the player

11.3 PolyXpress Test Code

References

- [1] “Ieee standard for software unit testing,” *ANSI/IEEE Std 1008-1987*, 1986.
- [2] “Systems and software engineering – vocabulary,” *ISO/IEC/IEEE 24765:2010(E)*, pp. 1–418, Dec 2010.
- [3] “Software and systems engineering software testing part 1:concepts and definitions,” *ISO/IEC/IEEE 29119-1:2013(E)*, pp. 1–64, Sept 2013.
- [4] K. Anantheswaran, “Istanbul.” [Online]. Available: <https://github.com/gotwarlost/istanbul>
- [5] A. Arkin, “Zombiejs.” [Online]. Available: <http://zombie.labnotes.org/>
- [6] “Grunt: The javascript task runner,” Bocoup. [Online]. Available: <http://gruntjs.com/>
- [7] “Chaijs,” chaijs. [Online]. Available: <http://chaijs.com/>
- [8] T. CI, “Travis.” [Online]. Available: <https://travis-ci.org/>
- [9] —, “Travis plans.” [Online]. Available: <https://travis-ci.com/plans>
- [10] D. Coates, “Node inspector.” [Online]. Available: <https://github.com/node-inspector/node-inspector>
- [11] W. W. W. Consortium, “Plan 2014,” 2012. [Online]. Available: <http://dev.w3.org/html5/decision-policy/html5-2014-plan.html>
- [12] “Polyxpress,” Expressive Technologies. [Online]. Available: <http://pxdev.cfapps.io/>
- [13] K. P. Frolin S. Ocasriza Jr. and B. Zorn, “Javascript errors in the wild: An empirical study,” 2011. [Online]. Available: http://ece.ubc.ca/~frolino/projects/jserr/tech_report.pdf
- [14] P. Heidegger and P. Thiemann, “Contract-driven testing of javascript code,” 2010. [Online]. Available: <https://proglang.informatik.uni-freiburg.de/jscontest/jscontest/>
- [15] M. Hevery, “jasmine-node.” [Online]. Available: <https://github.com/mhevery/jasmine-node>
- [16] A. Hidayat, “Phantomjs.” [Online]. Available: <http://phantomjs.org/>
- [17] E. Hieatt and R. Mee, “Going faster: Testing the web application,” *IEEE Software*, 2002. [Online]. Available: <http://www.ecs.csun.edu/~rlingard/COMP595VAV/GoingFaster.pdf>

- [18] Hyperakt and Vizzuality, “The evolution of the web.” [Online]. Available: <http://www.evolutionoftheweb.com/>
- [19] M. M. J. T. Javier Jess Gutierrez, Maria Jos Escalona, “Testing web application in practice,” in *First International Workshopl, WWV*, 2005. [Online]. Available: <http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.81.3910&rep=rep1&type=pdf#page=75>
- [20] “Jenkins organization,” Jenkins CI. [Online]. Available: <https://github.com/jenkinsci>
- [21] I. Joyent, “Node.js.” [Online]. Available: <http://nodejs.org/>
- [22] “Qunit,” The jQuery Foundation. [Online]. Available: <http://qunitjs.com/>
- [23] K. Kawaguchi, “Jenkins.” [Online]. Available: <http://jenkins-ci.org/>
- [24] V. Kongsli, “Security testing with selenium,” 2007. [Online]. Available: <http://www.kongsli.net/oopsla/OOPSLA'07%20Security%20Testing%20with%20Selenium%20v1.0.pdf>
- [25] “Magnumci,” Magnum CI. [Online]. Available: <https://magnum-ci.com>
- [26] A. Mardanov, “Todo app with express.js/node.js and mongodb.” [Online]. Available: <http://webapplog.com/todo-app-with-express-jsnode-js-and-mongodb/>
- [27] S. McConnell, “Code complete, 2nd edition,” 2004.
- [28] A. Mesbah and A. van Deursen, “Invariant-based automatic testing of ajax user interfaces,” 2009. [Online]. Available: <http://citeseerx.ist.psu.edu/viewdoc/download?doi=10.1.1.152.9899&rep=rep1&type=pdf>
- [29] S. Mirshokraie and A. Mesbah, “Jsart: Javascript assertion-based regression testing,” 2012. [Online]. Available: <http://www.ece.ubc.ca/~amesbah/docs/icwe12.pdf>
- [30] M. H. Mustafa Bozkurt and Y. Hassoun, “Testing web services: A survey,” King’s College London, Tech. Rep., 2010. [Online]. Available: <http://www.dcs.kcl.ac.uk/technical-reports/papers/TR-10-01.pdf>
- [31] D. A. Ostrowski, “Jaws: A javascript api for the efficient testing and integration of semantic web services,” 2007. [Online]. Available: <http://ftp.informatik.rwth-aachen.de/Publications/CEUR-WS/Vol-248/paper9.pdf>
- [32] A. B. Phillip Heidegger and P. Thiemann, “Dom transactions for testing javascript,” 2010. [Online]. Available: <https://proglang.informatik.uni-freiburg.de/jscontest/trans.pdf>

- [33] “Jasmine,” Pivotal Labs. [Online]. Available: <http://jasmine.github.io/>
- [34] H. Z. Qingning Huo and S. Greenwood, “A multi-agent software environment for testing web-based applications,” in *International Computer Software and Applications Conference*, 2003. [Online]. Available: <http://student.cse.fau.edu/~jsloan11/CEN6076/002.MasArch4Test.pdf>
- [35] B. Rady and R. Coffin, *Continuous Testing with Ruby, Rails, and JavaScript*. The Pragmatic Bookshelf, 2011, ch. Creating a JavaScript CT Environment.
- [36] “Semaphore,” RenderedText. [Online]. Available: <https://semaphoreapp.com/>
- [37] “Semaphore blog,” RenderedText. [Online]. Available: <https://semaphoreapp.com/blog/>
- [38] “Why amd?” RequireJS. [Online]. Available: <http://requirejs.org/docs/whyamd.html>
- [39] H. Reza and D. V. Gilst, “A framework for testing restful web services,” in *Seventh International Conference on Information Technology*, 2010. [Online]. Available: <http://ieeexplore.ieee.org/stamp/stamp.jsp?tp=&arnumber=5501468>
- [40] A. Rusu, “Nightwatchjs.” [Online]. Available: <http://nightwatchjs.org/>
- [41] B. Rydzewski and T. Burke, “Drone.io.” [Online]. Available: <https://drone.io/>
- [42] “Samsung rf4289hars,” Samsung. [Online]. Available: <http://www.samsung.com/us/appliances/refrigerators/RF4289HARS/XAA>
- [43] “Sauce labs,” Sauce Labs. [Online]. Available: <https://saucelabs.com/>
- [44] “Selenium,” SeleniumHQ. [Online]. Available: <http://docs.seleniumhq.org/>
- [45] J. D. F. T. D. D. A. P. Shay Artzi, Adam Kiezun and M. D. Ernst, “Finding bugs in dynamic web applications,” 2008. [Online]. Available: <http://dspace.mit.edu/bitstream/handle/1721.1/40249/MIT-CSAIL-TR-2008-006.pdf>
- [46] S. H. J. A. M. Shay Artzi, Julian Dolby and F. Tip, “A framework for automated testing of javascript web applications,” 2011. [Online]. Available: <https://cs.uwaterloo.ca/~ftip/pubs/icse2011artemis.pdf>
- [47] “intern,” SitePen Labs. [Online]. Available: theintern.io
- [48] “Browser usage statistics,” StatCounter. [Online]. Available: <http://gs.statcounter.com/#desktop-browser-ww-monthly-201402-201402-bar>
- [49] T. Tidwell, “Zombie phantom.” [Online]. Available: <https://github.com/travist/zombie-phantom>

- [50] P. Tonella and F. Ricca, “A 2-layer model for the white-box testing of web applications,” in *Telecommunications Energy Conference, 2004. INTELEC 2004. 26th Annual International*, 2004.
- [51] T. O. Valentin Dallmeier, Martin Burger and A. Zeller, “Webmate: A tool for testing web 2.0 applications,” 2012. [Online]. Available: <http://www.testfabrik.com/assets/pdf/webmate-jstools12.pdf>
- [52] “Mocha,” VisionMedia. [Online]. Available: <http://visionmedia.github.io/mocha/>
- [53] “Usage of javascript for websites,” W3Techs. [Online]. Available: <http://w3techs.com/technologies/details/cp-javascript/all/all>